Connect Jupyter Notebook with CDP Data Lake over AWS Athena

Data analysis and management using Amazon SageMaker

The final step is to make the AudienceDB and EventDB tables definitions available in a Jupyter notebook instance of Amazon SageMaker. Jupyter notebooks are popularly used among data scientists to visualize data, perform statistical analysis, do data manipulations, and make the data ready for machine learning work.

1. In the Amazon SageMaker console choose Create notebook instance.

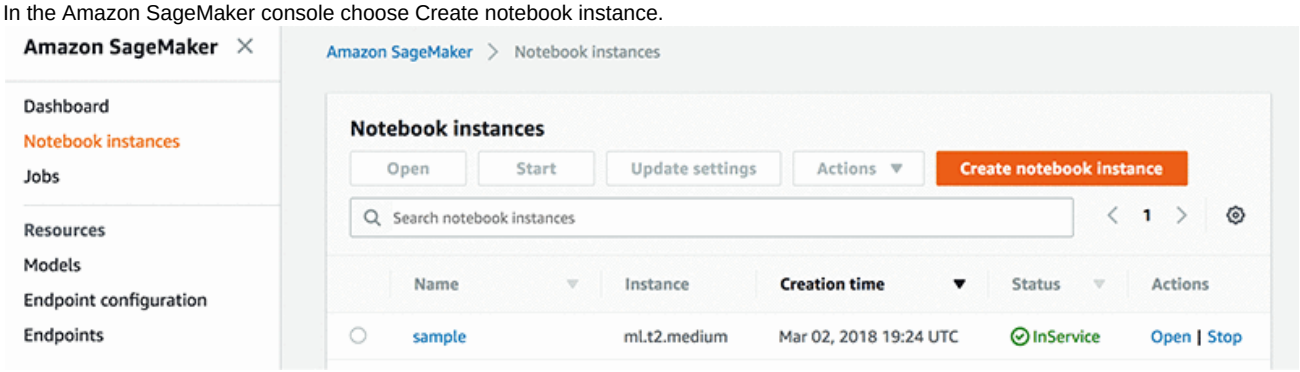

2. Under Notebook Instance settings populate the Notebook instance name, choose an instance type, and a role for the notebook instances in Amazon SageMaker to interact with Amazon S3. The SageMaker execution role needs to have necessary permission to Athena, the S3 buckets where the data resides, and KMS if encrypted.

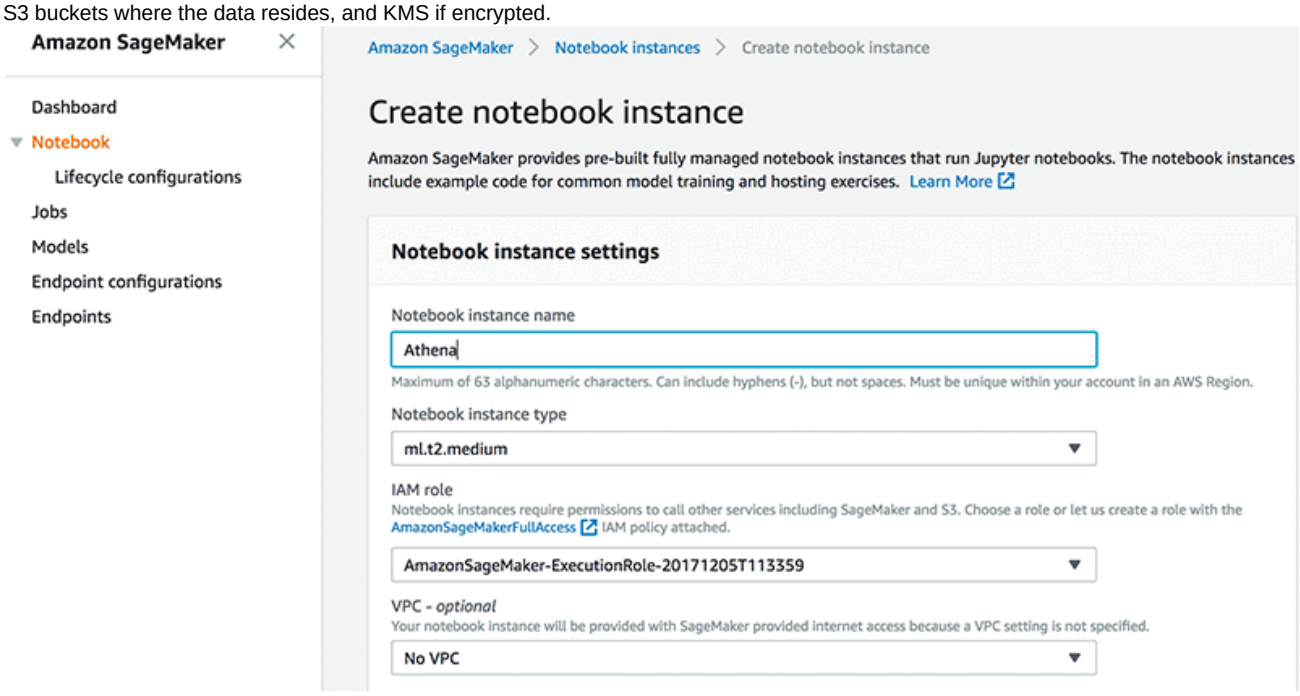

3. Wait for the Notebook instances to be created and the Status to change to InService.

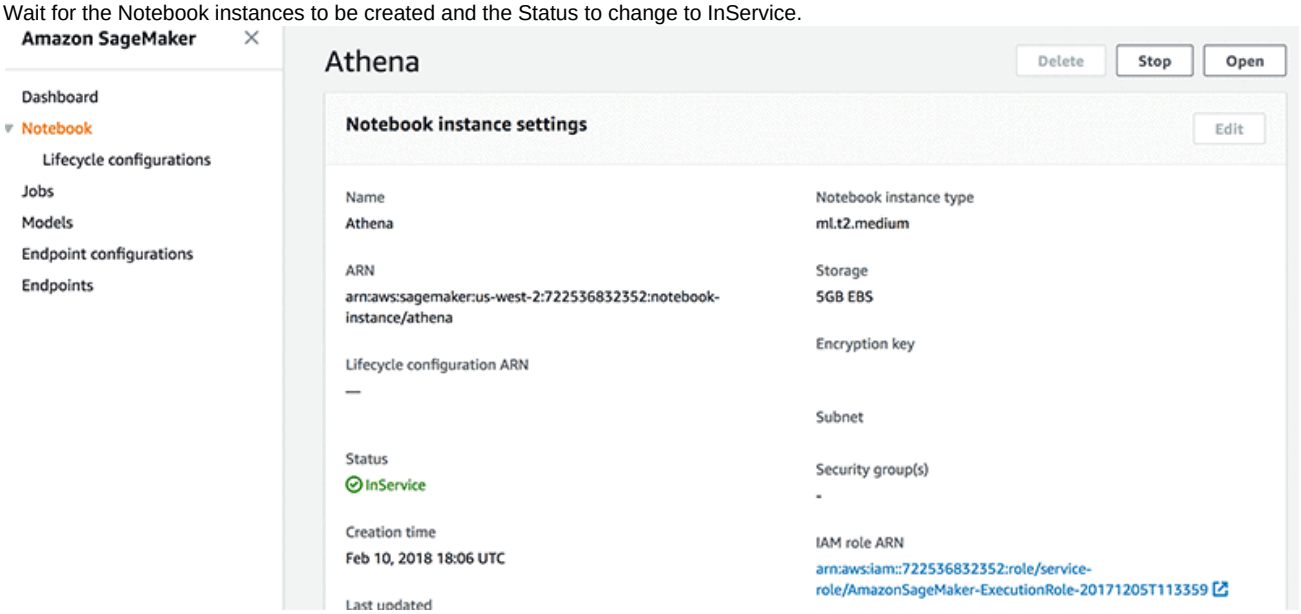

4. Choose the Open link, which will open the notebook interface in a separate browser window.

5. Click new to create a new notebook in Jupyter. Amazon SageMaker provides several kernels for Jupyter including support for Python 2 and 3, MXNet, TensorFlow, and PySpark. Choose Python as the kernel for this exercise as it comes with the Pandas library built in. Within the notebook, execute the following commands to install the Athena JDBC driver. PyAthena is a Python DB API 2.0 (PEP 249) co mpliant client for the Amazon Athena JDBC driver.

```
import sys
!{sys.executable} -m pip install PyAthena
```

```
In [1]: import sys
        !{sys.executable} -m pip install pyathena
```

```
Collecting pyathena
  Downloading https://files.pythonhosted.org/packages/26/97/a7fc04da461fb2f4b1cb5b886bbdfa38adc11f53218beb39d7c4564e5e0e/PyAthena-1.3.0-py2.py3-none-any.whl
Requirement already satisfied: boto3>=1.4.4 in /home/ec2-user/anaconda3/envs/python3/lib/python3.6/site-packages (from pyathena) (1.7.79)
Collecting future (from pyathena)
  Downloading https://files.pythonhosted.org/packages/00/2b/8d082ddfed935f3608cc61140df6dcbf0edea1bc3ab52fb6c29ae3e81e85/future-0.16.0.tar.gz (824kB)
    100% |████████████████████████████████| 829kB 4.0MB/s ta 0:00:01
Collecting tenacity>=4.1.0 (from pyathena)
  Downloading https://files.pythonhosted.org/packages/b5/02/f912867529807b879972d8000e23c2f67b8b3755171e1d3c2049e347a3c9/tenacity-5.0.2-py2.py3-none-any.whl
Requirement already satisfied: botocore>=1.5.52 in /home/ec2-user/anaconda3/envs/python3/lib/python3.6/site-packages (from pyathena) (1.10.79)
Requirement already satisfied: s3transfer<0.2.0,>=0.1.10 in /home/ec2-user/anaconda3/envs/python3/lib/python3.6/site-packages (from boto3>=1.4.4->pyathena) (0.1.13)
Requirement already satisfied: jmespath<1.0.0,>=0.7.1 in /home/ec2-user/anaconda3/envs/python3/lib/python3.6/site-packages (from boto3>=1.4.4->pyathena) (0.9.3)
Requirement already satisfied: six>=1.9.0 in /home/ec2-user/anaconda3/envs/python3/lib/python3.6/site-packages (from tenacity>=4.1.0->pyathena) (1.11.0)
Requirement already satisfied: docutils>=0.10 in /home/ec2-user/anaconda3/envs/python3/lib/python3.6/site-packages (from botocore>=1.5.52->pyathena) (0.14)
Requirement already satisfied: python-dateutil<3.0.0,>=2.1; python_version >= "2.7" in /home/ec2-user/anaconda3/envs/python3/lib/python3.6/site-packages (from botocore>=1.5.52->pyathena) (2.7.3)
Building wheels for collected packages: future
  Running setup.py bdist_wheel for future ... done
  Stored in directory: /home/ec2-user/.cache/pip/wheels/bf/c9/a3/c538d90ef17cf7823fa51fc701a7a7a910a80f6a405bf15b1a
Successfully built future
```

6. After the Athena driver is installed, you can use the JDBC connection to connect to Athena and populate the Pandas data frames. For data scientists, working with data is typically divided into multiple stages: munging and cleaning data, analyzing / modeling it, then organizing the results of the analysis into a form suitable for plotting or tabular display. Pandas is the ideal tool for all of these tasks.

```
from pyathena import connect
import pandas as pd
conn = connect(s3_staging_dir='<ATHENA QUERY RESULTS LOCATION>',
               region_name='<YOUR REGION, for example, us-west-2>')

df = pd.read_sql("SELECT * FROM athenaquery.<YOUR TABLE NAME> limit 8;", conn)
df
```

```
In [31]: from pyathena import connect
         import pandas as pd
         conn = connect(s3_staging_dir='s3://aws-athena-query-results-[redacted]/',
                        region_name='us-west-2')
         df = pd.read_sql("SELECT * FROM athenaquery.athenaquery_1518367613804 limit 10;", conn)
         df
```

Out[31]:

	year	month	dayofmonth	dayofweek	deptime	crsdeptime	arrtime	crsarrtime	uniquecarrier	flightnum	...	taxiin	taxiout	cancelled	cancellationcode	div
0	2008	8	29	5	2002	2002	2123	2134	OO	6650	...	6	14	0		
1	2008	8	29	5	2113	2105	2207	2155	OO	6651	...	4	15	0		
2	2008	8	29	5	1052	1059	1218	1232	OO	6653	...	7	21	0		
3	2008	8	29	5	1432	1437	1536	1539	OO	6655	...	10	9	0		
4	2008	8	29	5	1000	1003	1346	1344	OO	6656	...	4	28	0		
5	2008	8	29	5	2111	2110	2206	2149	OO	6659	...	7	9	0		
6	2008	8	29	5	1441	1443	1626	1638	OO	6661	...	4	10	0		
7	2008	8	29	5	1022	1022	1330	1332	OO	6662	...	4	18	0		
8	2008	8	29	5	1134	1136	1318	1318	OO	6663	...	4	19	0		
9	2008	8	29	5	1343	1349	1714	1722	OO	6663	...	12	8	0		

10 rows × 29 columns

Conclusion

The solution described in this blog post provides an automated way to catalog the incoming data as it comes into the data store, and it provides the ability to query the data for data manipulation and analysis. In addition, this scenario sets the stage for building more ML models through feature engineering, training, and scoring to gain more insights into your data and deliver significant business outcomes.